

# Speeding up Python using Cython

Rolf Boomgaarden · Thiemo Gries · Florian Letsch

Universität Hamburg

November 28th, 2013

# What is Cython?

- ▶ Compiler, compiles Python-like code to C-code
- ▶ Code is still executed in the Python runtime environment
- ▶ But is compiled to native machine code instead of Python bytecode
- ▶ Can result in more speed and easy wrapping of C libraries

# Cython workflow

## first approach

- ▶ Write a helloworld.pyx source file

```
1 print "Hello_World"
```

- ▶ Run the Cython compiler to generate a C file

```
1 $ cython helloworld.pyx
```

- ▶ Run a C compiler to generate a compiled library

```
1 $ gcc [...] -o helloworld.so helloworld.c
```

- ▶ Run the Python interpreter and import the module

```
1 >>> import helloworld
2 Hello World
```

# Cython workflow

## first approach

- ▶ Write a helloworld.pyx source file

```
1 print "Hello_World"
```

- ▶ Run the Cython compiler to generate a C file

```
1 $ cython helloworld.pyx
```

- ▶ Run a C compiler to generate a compiled library

```
1 $ gcc [...] -o helloworld.so helloworld.c
```

- ▶ Run the Python interpreter and import the module

```
1 >>> import helloworld
2 Hello World
```

# Cython workflow

## second approach

- ▶ Write a helloworld.pyx source file
- ▶ Write a setup.py file with compile information

```
1 from distutils.core import setup
2 from Cython.Build import cythonize
3
4 setup(
5     ext_modules = cythonize("helloworld.pyx")
6 )
```

- ▶ Let Python compile the file

```
1 $ python setup.py build_ext --inplace
```

- ▶ Run the Python interpreter and import the module

# Cython workflow

## second approach

- ▶ Write a helloworld.pyx source file
- ▶ Write a setup.py file with compile information

```
1 from distutils.core import setup
2 from Cython.Build import cythonize
3
4 setup(
5     ext_modules = cythonize("helloworld.pyx")
6 )
```

- ▶ Let Python compile the file

```
1 $ python setup.py build_ext --inplace
```

- ▶ Run the Python interpreter and import the module

# Cython workflow

third approach: pyximport

- ▶ Write a helloworld.pyx source file
- ▶ Use Pyximport

```
1 >>> import pyximport; pyximport.install()  
2 >>> import helloworld  
3 Hello World
```

**KOGS**



**WINTER**

**of**

**CODE**

**FIRST**

**2014**

**YEAR**



# Pure Python

```
1 import numpy as np
2
3 def my_add(a, b):
4     (...) # validate parameter
5
6     dtype = a.dtype
7     height = a.shape[0]
8     width = a.shape[1]
9
10    result = np.zeros((height, width), dtype=dtype)
11
12    for y in range(height):
13        for x in range(width):
14            result[y,x] = a[y,x] + b[y,x]
15
16    return result
```

Listing 1: add1.py

# Pure Python

```
1 import numpy as np
2
3 def my_add(a, b):
4     (...) # validate parameter
5
6     dtype = a.dtype
7     height = a.shape[0]
8     width = a.shape[1]
9
10    result = np.zeros((height, width), dtype=dtype)
11
12    for y in range(height):
13        for x in range(width):
14            result[y,x] = a[y,x] + b[y,x]
15
16    return result
```

Listing 1: add1.py

Time: ~19 minutes (2048x2048, 100x)

# Python run through Cython

```
1 import numpy as np
2
3 def my_add(a, b):
4     (...) # validate parameter
5
6     dtype = a.dtype
7     height = a.shape[0]
8     width = a.shape[1]
9
10    result = np.zeros((height, width), dtype=dtype)
11
12    for y in range(height):
13        for x in range(width):
14            result[y,x] = a[y,x] + b[y,x]
15
16    return result
```

Listing 2: add2.pyx

# Python run through Cython

```
1 import numpy as np
2
3 def my_add(a, b):
4     (...) # validate parameter
5
6     dtype = a.dtype
7     height = a.shape[0]
8     width = a.shape[1]
9
10    result = np.zeros((height, width), dtype=dtype)
11
12    for y in range(height):
13        for x in range(width):
14            result[y,x] = a[y,x] + b[y,x]
15
16    return result
```

Listing 2: add2.pyx

Time: ~16 minutes (2048x2048, 100x)

# Cython: Adding types

```
1 import numpy as np
2 cimport numpy as np
3
4 DTYPE = np.uint8
5 ctypedef np.uint8_t DTYPE_t
6
7 def my_add(np.ndarray a, np.ndarray b):
8     (...) # validate parameter
9
10    cdef int height = a.shape[0]
11    cdef int width = a.shape[1]
12    cdef np.ndarray result = np.zeros((height, width), dtype=DTYPE)
13
14    cdef int x, y
15    for y in range(height):
16        for x in range(width):
17            result[y,x] = a[y,x] + b[y,x]
18
19    return result
```

Listing 3: add3.pyx

# Cython: Adding types

```
1 import numpy as np
2 cimport numpy as np
3
4 DTYPE = np.uint8
5 ctypedef np.uint8_t DTYPE_t
6
7 def my_add(np.ndarray a, np.ndarray b):
8     (...) # validate parameter
9
10    cdef int height = a.shape[0]
11    cdef int width = a.shape[1]
12    cdef np.ndarray result = np.zeros((height, width), dtype=DTYPE)
13
14    cdef int x, y
15    for y in range(height):
16        for x in range(width):
17            result[y,x] = a[y,x] + b[y,x]
18
19    return result
```

Listing 3: add3.pyx

Time: ~16 minutes (2048x2048, 100x)

# Cython: Efficient indexing

```
1 import numpy as np
2 cimport numpy as np
3
4 DTYPE = np.uint8
5 ctypedef np.uint8_t DTYPE_t
6
7 def my_add(np.ndarray[DTYPE_t, ndim=2] a, np.ndarray[DTYPE_t, ndim=2] b
8     ):
9     (...) # validate parameter
10
11     cdef int height = a.shape[0]
12     cdef int width = a.shape[1]
13     cdef np.ndarray[DTYPE_t, ndim=2] result = np.zeros((height, width),
14         dtype=DTYPE)
15
16     cdef int x, y
17     for y in range(height):
18         for x in range(width):
19             result[y,x] = a[y,x] + b[y,x]
20
21     return result
```

Listing 4: add4.pyx

# Cython: Efficient indexing

```
1  import numpy as np
2  cimport numpy as np
3
4  DTYPE = np.uint8
5  ctypedef np.uint8_t DTYPE_t
6
7  def my_add(np.ndarray[DTYPE_t, ndim=2] a, np.ndarray[DTYPE_t, ndim=2] b
8      ):
9      (...) # validate parameter
10
11      cdef int height = a.shape[0]
12      cdef int width = a.shape[1]
13      cdef np.ndarray[DTYPE_t, ndim=2] result = np.zeros((height, width),
14          dtype=DTYPE)
15
16      cdef int x, y
17      for y in range(height):
18          for x in range(width):
19              result[y,x] = a[y,x] + b[y,x]
```

Listing 4: add4.pyx

Time: 1.249 seconds (2048x2048, 100x)



# Cython: Don't check boundaries

```
1 import numpy as np; cimport numpy as np; cimport cython
2
3 DTYPE = np.uint8
4 ctypedef np.uint8_t DTYPE_t
5
6 @cython.boundscheck(False)
7 def my_add(np.ndarray[DTYPE_t, ndim=2] a, np.ndarray[DTYPE_t, ndim=2] b
8     ):
9     (...) # validate parameter
10
11     cdef int height = a.shape[0]
12     cdef int width = a.shape[1]
13     cdef np.ndarray[DTYPE_t, ndim=2] result = np.zeros((height, width),
14         dtype=DTYPE)
15
16     cdef int x, y
17     for y in range(height):
18         for x in range(width):
19             result[y,x] = a[y,x] + b[y,x]
20
21     return result
```

Listing 5: add5.pyx


# Cython: Don't check boundaries

```
1 import numpy as np; cimport numpy as np; cimport cython
2
3 DTYPE = np.uint8
4 ctypedef np.uint8_t DTYPE_t
5
6 @cython.boundscheck(False)
7 def my_add(np.ndarray[DTYPE_t, ndim=2] a, np.ndarray[DTYPE_t, ndim=2] b
8     ):
9     (...) # validate parameter
10
11     cdef int height = a.shape[0]
12     cdef int width = a.shape[1]
13     cdef np.ndarray[DTYPE_t, ndim=2] result = np.zeros((height, width),
14         dtype=DTYPE)
15
16     cdef int x, y
17     for y in range(height):
18         for x in range(width):
19             result[y,x] = a[y,x] + b[y,x]
20
21     return result
```


Listing 5: add5.pyx

Time: 1.086 seconds (2048x2048, 100x)

# Conclusion



# Conclusion



# FIXME Flo

FIXME: speedup, how much fun to work with, easy use of features, documentation etc

# The end

EOF

(FIXME: better joke)